



# Practical Approaches to Optimizing Virtual Memory Buffer Management

第七組 鄧雅文、王偉力、吳吉加、周哲瑋



# outline

1.Motivation

2.Introduction

3.Method

4.Result

5.Conclusion



# Motivation

## 1. Rapid Growth of AI and Big Data

- Exponential increase in data volume, heightened demand for data access efficiency

## 2. Optimizing System from Middleware Perspective

- Not just model optimization, but also enhancing hardware-software interaction speed

## 3. Improving Database Performance

- Based on VMCACHE code, exploring software-level improvements for better data access efficiency

## 4. Goals

- Provide efficient data management solutions for applications like AI or big data



# Introduction

## Virtual-Memory Assisted Buffer Management

Preprint accepted for publication at SIGMOD 2023

Viktor Leis

Technische Universität München  
leis@in.tum.de

Adnan Alhomssi

Friedrich-Alexander-Universität  
Erlangen-Nürnberg  
adnan.alhomssi@fau.de

Tobias Ziegler

Technische Universität Darmstadt  
tobias.ziegler@cs.tu-darmstadt.de

Yannick Loeck

Technische Universität Hamburg  
yannick.loeck@tuhh.de

Christian Dietrich

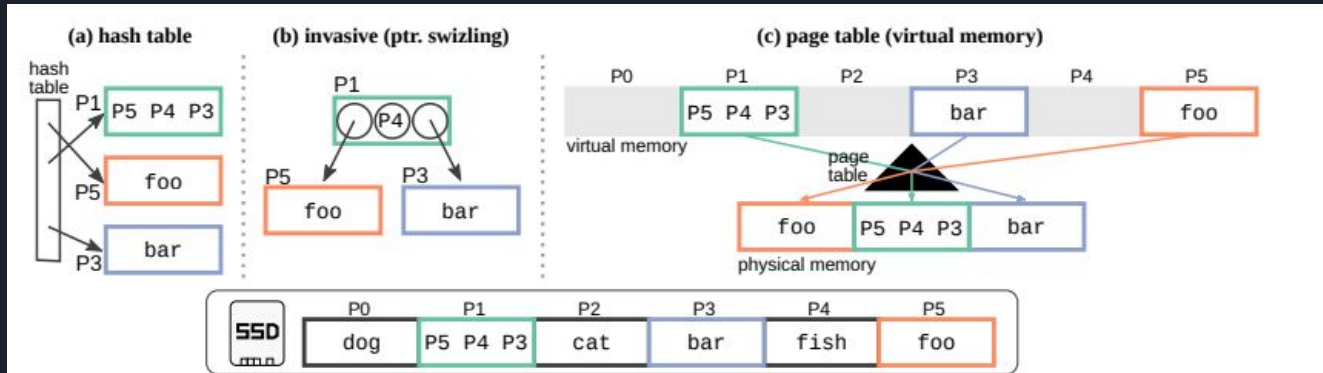
Technische Universität Hamburg  
christian.dietrich@tuhh.de

# Introduction

Larger-than-memory Database

VMCACHE : VM + PT + TLB, DBMS-driven , easy to implement

page management, page states, page table length





# Method

## Improvements Based on Original Code







Five methods tested with random access and TPC-C benchmark

1. **Mutex:** Addresses synchronization issues
2. **Futex:** Handles synchronization with a fast userspace mutex, lighter and faster than mutex
3. **Merge:** B-tree inner node merge
4. **Yield-Opt:** Dynamically adjusts waiting strategies
5. **Bitmap:** Manages free space

**Implementation Division:** method 1~4 周哲瑋, 5 鄧雅文

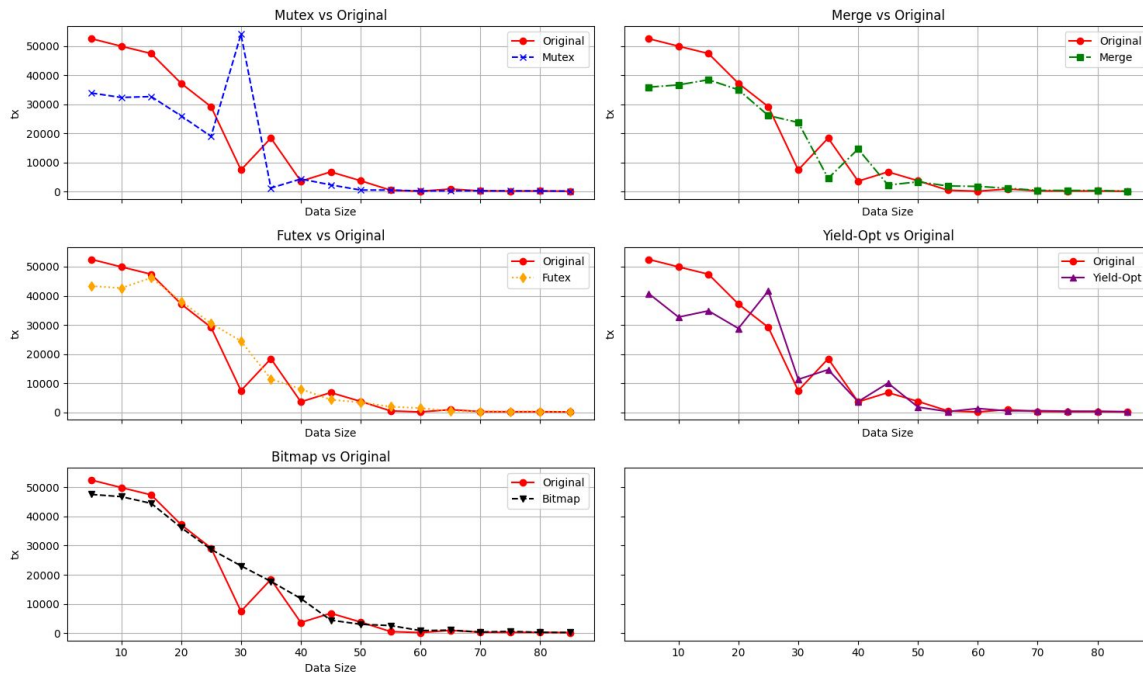
# Result

Random Lookup: tx



# Result

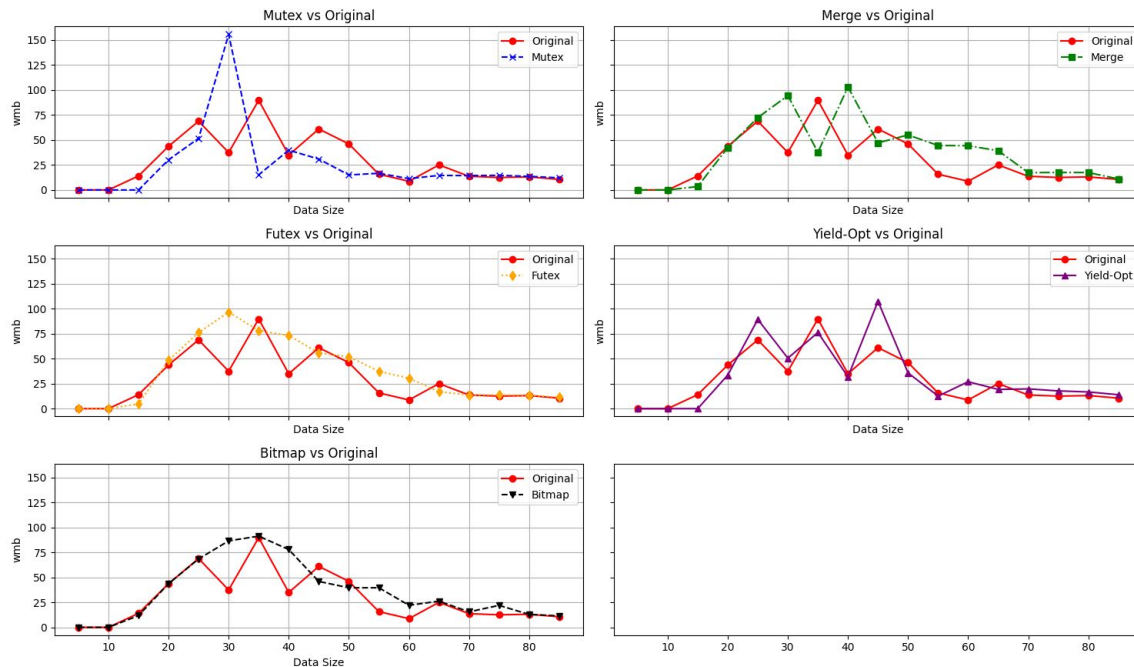
TPC-C: tx





# Result

TPC-C: wmb





# Conclusion

Implementation of 5 Improvement Methods:

Overall Performance:

- Futex, Yield-Opt, and Bitmap show slight advantages over the original version

Testing Environment:

- Used Google Colab, limited to testing with up to 4 threads
- Unable to test with varying thread counts
- Long testing times

More Details : <https://github.com/5000user5000/vmcache/tree/master>



# Thanks for listening